


Class 07

Tim

Before we get into clustering methods let's make some sample data to cluster where we know what the answer should be

To help with this I will use the `rnorm()` function

```
hist(rnorm(150000, mean = c(-3,3)))
```



```
n = 30  
c(rnorm(n, mean = 3), rnorm(n, mean = -3))
```

```
[1] 3.0673744 3.3921853 4.9633241 2.5177356 2.6219122 2.9027532  
[7] 2.8145837 2.4239096 3.1507842 2.7433154 3.0907909 3.7790824
```

```

[13]  1.1990955  1.5348185  4.4188325  3.1053706  2.5333590  2.0666045
[19]  4.0228372  1.9275282  3.7377766  3.7639632  3.5439448  4.0378093
[25]  2.3301366  1.8500144  2.3368215  3.0600333  4.1878627  4.5378783
[31] -1.8516578 -2.8796135 -3.9116693 -4.6042732 -0.7613992 -2.6786801
[37] -5.0145288 -3.5069505 -3.7206075 -3.5670613 -3.5723703 -2.1933569
[43] -4.6943384 -1.1339635 -4.0082621 -4.8991815 -4.1265986 -3.0615853
[49] -3.1763520 -3.6077591 -4.0878717 -1.8227651 -0.5573039 -2.5862250
[55] -3.1308867 -4.4812359 -3.5378964 -3.7234471 -2.0955326 -3.5964401

```

```

n = 30
x <- c(rnorm(n, mean = 3), rnorm(n, mean = -3))
y <- rev(x)

z <- cbind(x,y)
z

```


```

      x      y
[1,]  1.453668 -3.217778
[2,]  5.101532 -3.664622
[3,]  2.629980 -2.455965
[4,]  3.712873 -3.245920
[5,]  3.418594 -2.030575
[6,]  2.853070 -1.806018
[7,]  4.145809 -3.594470
[8,]  4.263497 -3.634083
[9,]  4.168928 -1.251198
[10,]  3.158374 -4.523698
[11,]  1.016139 -2.774650
[12,]  2.837173 -4.062118
[13,]  3.251956 -1.817334
[14,]  2.420578 -3.359221
[15,]  2.562483 -3.650044
[16,]  4.736963 -1.127376
[17,]  2.802062 -3.572527
[18,]  4.059955 -2.707002
[19,]  2.807184 -1.933435
[20,]  2.273081 -3.534090
[21,]  3.545364 -3.350019
[22,]  2.301536 -1.886402
[23,]  2.671902 -1.803242
[24,]  3.691481 -3.431428
[25,]  1.701671 -1.651517

```

```
[26,] 2.291465 -1.900011
[27,] 2.642904 -2.235410
[28,] 2.222915 -3.696900
[29,] 1.788067 -2.693113
[30,] 2.535225 -3.183098
[31,] -3.183098 2.535225
[32,] -2.693113 1.788067
[33,] -3.696900 2.222915
[34,] -2.235410 2.642904
[35,] -1.900011 2.291465
[36,] -1.651517 1.701671
[37,] -3.431428 3.691481
[38,] -1.803242 2.671902
[39,] -1.886402 2.301536
[40,] -3.350019 3.545364
[41,] -3.534090 2.273081
[42,] -1.933435 2.807184
[43,] -2.707002 4.059955
[44,] -3.572527 2.802062
[45,] -1.127376 4.736963
[46,] -3.650044 2.562483
[47,] -3.359221 2.420578
[48,] -1.817334 3.251956
[49,] -4.062118 2.837173
[50,] -2.774650 1.016139
[51,] -4.523698 3.158374
[52,] -1.251198 4.168928
[53,] -3.634083 4.263497
[54,] -3.594470 4.145809
[55,] -1.806018 2.853070
[56,] -2.030575 3.418594
[57,] -3.245920 3.712873
[58,] -2.455965 2.629980
[59,] -3.664622 5.101532
[60,] -3.217778 1.453668
```

```
plot(z)
```



K-means clustering

the function in base R for kmeans


```
km <- kmeans(z, centers = 2)
```

Q. print out the cluster membership vector (i.e. our main answer)


```
km$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
plot(z, col=c("red","blue"))
```




```
plot(x, col=c(1,2))
```




plot with clustering center

```
plot(z, col= km$cluster)
points(km$centers, col = "blue", pch=15, cex=2)
```



Q. can you cluster our data in z into 4 clusters?

```
km4 <- kmeans(z, centers = 4)
plot(z, col=km4$cluster)
points(km4$centers, col ="blue", pch=15, cex=2)
```



Hierarchical Clustering

the main function for hierarchical Clustering in base R is called `hclust`

unlike `kmeans()` I can not just pass in my data as input I first need a distance matrix from data.

```
d <- dist(z)
hc <- hclust(d)
hc
```

Call:


```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

there is a specific `hclust plot()` method..

```
plot(hc)
abline(h=10, col="red")
```

Cluster Dendrogram



the longer the cluster bar the larger the distance and the more space to score a difference between the clusters

to get my main clustering results (i.e. the membership vector) can “cut” my tree at a given height. To add this I will use the `cutree`

```
grps <- cutree(hc, h=10)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Principal Component Analysis

reduce the features dimensionality while only losing a small amount of info (follows the best fit line through points)


```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

```
pairs(x, col= rainbow(10), pch=16)
```



the main function to do PCA in base R is called `prcomp()`

`t()` function is transpose to switch the x axis to make it y axis

`t(x)`

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175
Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139
	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes		
England	720	253	488		198	
Wales	874	265	570		203	
Scotland	566	171	418		220	
N.Ireland	1033	143	355		187	
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks	
England	360	1102	1472	57	1374	
Wales	365	1137	1582	73	1256	
Scotland	337	957	1462	53	1572	
N.Ireland	334	674	1494	47	1506	
	Alcoholic_drinks	Confectionery				
England	375	54				

Wales	475	64
Scotland	458	62
N.Ireland	135	41

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

#Proportion variance is how much each PC captured (i.e., PC1 captured 67% of the data 0.67)

#cumulative proportion you add PC1 to PC2 ($0.67 + 0.29$)

let's see what is inside our result object `pca` that we just calculated

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class

```
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13